

## ST3189 Coursework codes

200664383

2023-04-02

<https://www.openml.org/data/download/1586225/php0iVrYT>

<https://www.openml.org/data/download/21377442/file16a868cf35f5.arff>

#Regression-----

#CART on runtime using predictor variables

#install and load 'rpart' package which provides the functions for building decision trees  
(for CART) install.packages("rpart") install.packages("lattice") install.packages("caret")  
library(ggplot2) library(rpart) library(caret) library(readr) library(lattice)

#set working directory to import dataset setwd("C:/Users/User/Documents/ST3189  
Machine Learning/ST3189 coursework")

#load dataset CPMP <- read.csv("CPMP.csv")

#remove duplicates CPMP\_nd <- unique(CPMP)

### Set random seed to make it reproducible

set.seed(123)

### Split data into train set (70%) and test set (30%)

train\_indices <- sample(nrow(CPMP\_nd), round(0.7 \* nrow(CPMP\_nd))) CPMP\_train\_data  
<- CPMP\_nd[train\_indices, ] CPMP\_test\_data <- CPMP\_nd[-train\_indices, ]

#check class of train set class(CPMP\_train\_data\$runtime)

#model the training set (CART) model\_CPMP <- rpart(runtime ~ stacks + tiers +  
overstowing.stack.pct + empty.stack.pct + left.density, data = CPMP\_train\_data, method =  
"anova")

#print the summary of the model summary(model\_CPMP)

#plot the tree par(mar = c(5, 5, 5, 5)) #set size of plot area windows(width=10, height=12)  
plot(model\_CPMP) text(model\_CPMP, use.n=TRUE, all=TRUE, cex=0.8)

```
#use test set to make predictions pred_CPMP <- predict(model_CPMP, newdata =
CPMP_test_data)
```

## Evaluate model using MSE, R-squared, and MAE

```
mse <- mean((CPMP_test_data$runtime - pred_CPMP)^2) r_squared <-
1 - cor(CPMP_test_data$runtime, pred_CPMP)^2 mae <- mean(abs(CPMP_test_data$runtime -
pred_CPMP))
```

## Print evaluation metrics

```
cat("Mean Squared Error (MSE):", mse, "\n") cat("R-squared:", r_squared, "\n") cat("Mean
Absolute Error (MAE):", mae, "\n")
```

#since the r squared value is near 0 and the mse and mae is very high , we will run PCA to reduce the complexity of the dataset, before modelling it again with CART

```
#PCA-----
#load required libraries library(caret) library(randomForest) library(dplyr) library(tidyr)
```

## Select the predictor variables

```
predictors <- c("stacks", "tiers", "overstowing.stack.pct", "empty.stack.pct", "left.density")
```

## Scale and perform PCA on the training data

```
CPMP_scaled_train <- scale(CPMP_train_data[, predictors])

pca <- prcomp(CPMP_scaled_train[, predictors]) cpmp_train_pca <- predict(pca,
CPMP_scaled_train)

plot(pca)

#choose principal components to train model cpmp_train_model <- data.frame(PC1 =
cpmp_train_pca[,1], PC2 = cpmp_train_pca[,2], PC3 = cpmp_train_pca[,3], runtime =
CPMP_train_data$runtime)
```

## Remove all rows with missing values from dataset

```
cpmp_train_model <- na.omit(cpmp_train_model)
```

## Train a CART model on the PCA-transformed data

```
cpmp_cart_model <- train(runtime ~ ., data = cpmp_train_model, method = "rpart",  
trControl = trainControl(method = "cv", number = 10))
```

## Print a summary of model

```
print(cpmp_cart_model)  
  
#predict runtime with train set cpmp_train_predicted <- predict(cpmp_cart_model,  
newdata = cpmp_train_model)
```

## Scale testing data using the same standard deviation and as the train data

```
CPMP_scaled_test <- scale(CPMP_test_data[, predictors], center = attr(CPMP_scaled_train,  
"scaled:center"), scale = attr(CPMP_scaled_train, "scaled:scale"))
```

## Project test data onto the principal components obtained from training data

```
cpmp_test_pca <- predict(pca, CPMP_scaled_test)
```

## View the first few rows of the test data after PCA

```
head(cpmp_test_pca)  
  
cpmp_test_model <- data.frame(PC1 = cpmp_test_pca[,1], PC2 = cpmp_test_pca[,2], PC3 =  
cpmp_test_pca[,3], runtime = CPMP_test_data$runtime)  
  
#Predict runtime with scaled test set cpmp_test_predicted <- predict(cpmp_cart_model,  
newdata = cpmp_test_model)
```

## Print the summary of model

```
print(cpmp_test_predicted)
```

## Calculate RMSE for train and test predictions

```
train_rmse <- RMSE(cpmp_train_predicted, CPMP_train_data$runtime)  
test_rmse <- RMSE(cpmp_test_predicted, CPMP_test_data$runtime)
```

## Print RMSE values

```
cat(paste("Train RMSE:", train_rmse, "")) cat(paste("Test RMSE:", test_rmse, ""))  
install.packages("rpart.plot") library(rpart.plot)
```

## Plot actual vs predicted runtime for train data

```
plot(CPMP_train_data$runtime, cmpmp_train_predicted, main = "Actual vs Predicted  
Runtime (Train Data)", xlab = "Actual Runtime", ylab = "Predicted Runtime")
```

## Plot actual vs predicted runtime for test data

```
plot(CPMP_test_data$runtime, cmpmp_test_predicted, main = "Actual vs Predicted Runtime  
(Test Data)", xlab = "Actual Runtime", ylab = "Predicted Runtime")
```

```
#Evaluate the performance of the CART model using mean squared error (MSE) and R-  
squared mse <- mean((cmpmp_test_predicted - CPMP_test_dataruntime)^2) rsq <-  
-cor(cmpmp_test_predicted, CPMP_test_dataruntime)^2
```

```
#pruning the decision tree
```

```
#showing which variable has the strongest relevance to determining runtime  
summary(pca) print(pca$rotation[,1])
```

```
#Since CART modelling still produced a low R-squared value and high MSE value, we will  
try randomForest to model the dataset
```

## RandomForest

---

```
#Fit a random forest model using randomForest() function rf_model <-  
randomForest(runtime ~ ., data = CPMP_train_data, ntree = 500, importance = TRUE)
```

```
#Use the predict() function to make predictions on the test set using fitted model  
prediction_CPMP <- predict(rf_model, newdata = CPMP_test_data)
```

```
#Evaluate performance of the model using mean squared error (MSE) and R-squared  
mse2 <- mean((predictions_CPMP - CPMP_test_dataruntime)^2)
```

```
rsq2 <- -cor(predictions_CPMP, CPMP_test_dataruntime)^2
```

```
#Plotting the out-of-bag error rate versus the number of trees in forest using the plot()  
function plot(rf_model)
```

```
#Visualize the predicted versus observed values of the target variable using a scatter plot  
prediction_CPMP <- predict(rf_model, CPMP_test_data) windows(width=10, height=8)
```

```
# adjusting width and height of plot (not sure if want to keep this) plot(predictions_CPMP,
CPMP_test_data$runtime, main = "Predicted vs Observed Values") abline(0, 1, col = "red")

#randomForests on overstowage.pct

rf_model2 <- randomForest(overstowage.pct ~ ., data = CPMP_train_data, ntree = 500,
importance = TRUE)

#Use the predict() function to make predictions on the test set using the fitted model
prediction_CPMP2 <- predict(rf_model2, newdata = CPMP_test_data)

#Evaluate the performance of the model using mean squared error (MSE) and R-squared
mse2 <- mean((prediction_CPMP2 - CPMP_test_data$overstowage.pct)^2)

rsq2 <- 1 - cor(prediction_CPMP2, CPMP_test_data$overstowage.pct)^2

#Plot the out-of-bag error rate versus the number of trees in the forest using the plot()
function plot(rf_model2) {
  plot(rf_model2)
}

#linear regression on overstowage
```

## Define the predictor variables

```
predictors <- c("stacks", "tiers", "left.density")
```

## Fit a linear regression model to predict overstacking.pct

```
model <- lm(overstowage.pct ~ stacks + tiers + left.density, data = CPMP_train_data)

#summarize the model summary(model)
```

## Predict overstowage.pct values using the model and predictor variables

```
prediction <- predict(model, newdata = CPMP_test_data[, predictors])

#visualise results library(ggplot2) library(dplyr) ggplot(CPMP_test_data, aes(x =
prediction, y = overstowage.pct)) + geom_point() + geom_smooth(method = "lm", se =
FALSE) + labs(title = "Linear Regression Prediction vs Actual Overstowage Percentage", x =
"Predicted Overstowage Percentage", y = "Actual Overstowage Percentage")

#Classification
```

```
#install and load necessary libraries install.packages("randomForest")
install.packages("ggplot2") install.packages("lattice") library(rpart) library(randomForest)
library(ggplot2) library(caret) library(readr) library(lattice)

#set working directory to import dataset setwd("C:/Users/User/Documents/ST3189
Machine Learning/ST3189 coursework")

#load data bd <- read.csv("blood_donation.csv.csv")

#remove duplicates bd_nd <- unique(bd)
```

## Set random seed to make it reproducible

```
set.seed(123)
```

## Split data into train set (70%) and test set (30%)

```
train_indices <- sample(nrow(bd_nd), round(0.7 * nrow(bd_nd))) bd_train_data <-
bd_nd[train_indices, ] bd_test_data <- bd_nd[-train_indices, ]

#check class of bd train set class(bd_train_data$Class)

#model the training set (CART) model <- rpart(Class ~ recency + frequency, data =
bd_train_data)

#print the summary of the model summary(model)

#plot the tree par(mar = c(5, 5, 5, 5)) #set size of plot area windows(width=10, height=8)
plot(model) text(model, use.n=TRUE, all=TRUE, cex=0.8)

#use test set to evaluate performance

pred <- predict(model, newdata = bd_test_data)

#compare predicted value to actual value table(pred, bd_test_data$Class)

accuracy <- sum(pred == bd_test_data$Class) / length(bd_test_data$Class)

precision <- sum(pred & bd_test_data$Class) / sum(pred)

recall <- sum(pred & bd_test_data$Class)/sum(bd_test_data$Class)

f1_score <- 2 * precision * recall / (precision + recall)

#Load the necessary libraries

library(dplyr) library(tidyr) library(ggplot2) library(caret)
```

## Check if the target variable is binary

```
if (length(unique(bd_nd$Class)) > 2) { stop("Target variable is not binary") }
```

## Replace 1 with 0 and 2 with 1 in the Class column

```
bdClass <- ifelse(bdClass == 1, 0, 1) bdClass <- ifelse(bdClass == 2, 1, bd$Class)
```

## Save the modified dataset

```
write.csv(bd, "bd.csv", row.names = FALSE)
```

## Set random seed to make it reproducible

```
set.seed(123)
```

## Split data into train set (70%) and test set (30%)

```
train_indices <- sample(nrow(bd), round(0.7 * nrow(bd))) bd_train_data <-  
bd[train_indices, ] bd_test_data <- bd[-train_indices, ]
```

## Perform logistic regression

```
logistic_model <- glm(Class ~ recency + frequency, data = bd_train_data, family =  
"binomial")
```

## Making predictions on the test data

```
pred <- predict(logistic_model, newdata = bd_test_data, type = "response") predicted_class  
<- ifelse(pred > 0.5, 1, 0)
```

```
predicted_class <- factor(predicted_class) bd_test_data$Class <- factor(bd_test_data$Class)
```

## Check the levels of both variables

```
levels(predicted_class) levels(bd_test_data$Class)
```

## Re-level the predicted variable to match the reference variable

```
predicted_class <- factor(predicted_class, levels = levels(bd_test_data$Class))
```

## Check the levels again to make sure they match

```
levels(predicted_class) levels(bd_test_data$Class)

#Evaluate the model confusion_matrix <- confusionMatrix(data = predicted_class,
reference = bd_test_data$Class)

#visualizing the confusion matrix mosaicplot(confusion_matrix$table, main = "Confusion
Matrix")
```

## randomForest-----

```
library(randomForest)
```

## Set seed to make it reproducible

```
set.seed(123)

bd <- read.csv("bd.csv") bd_nd <- na.omit(bd_nd)

#Split the data into training set (70%) and testing set (30%)

train_row <- sample(nrow(bd), round(0.7 * nrow(bd)))

bd_train_data <- bd_nd[train_row, ]

bd_test_data <- bd_nd[-train_row, ]
```

## Train randomForest model

```
bd_model_rf <- randomForest(Class ~ ., data = bd_train_data, ntree = 80, mtry =
sqrt(ncol(bd_train_data)))
```

## Make predictions on test data

```
pred_bd <- predict(bd_model_rf, bd_test_data)
```

## Evaluate accuracy of predictions

```
accuracy <- mean(pred_bd == bd_test_data$Class)
```

## Re-level the predicted variable to match the reference variable

```
pred_bd <- factor(pred_bd, levels = levels(bd_test_data$Class))
```



response has 5 or fewer unique values, indicating that randomForest is not suitable to model this dataset——